
Supplementary Material for Masked Autoencoder with Time Series Generation

Anonymous Author(s)

Affiliation

Address

email

- 1 The source code is under the folder *MAI*. The reader may also download it from Github. You can
2 check the results in jupyter notes *ExtraMAE.ipynb*, *MaskRatio.ipynb*, and *Imputation.ipynb* to save
3 your time. The jupyter notes only show final results, and readers have to download the full datasets
4 and model files to run the jupyter notes. Since the full datasets and model files are too large, the
5 supporting codes and files for *ExtraMAE.ipynb*, *MaskRatio.ipynb*, and *Imputation.ipynb* are available
6 upon request.

7 **1 Visualization**

- 8 We provide t-SNE [1] plots for all candidate models (ExtraMAE in row1, TimeGAN [2] in row2,
9 RC-GAN [3] in row3, and C-RNN-GAN [4] in row4). We randomly select six features to plot.
10 Each column visualizes one feature. Features from left to right are Close, Volume, Sine1, Sine4,
11 Application, and T1.

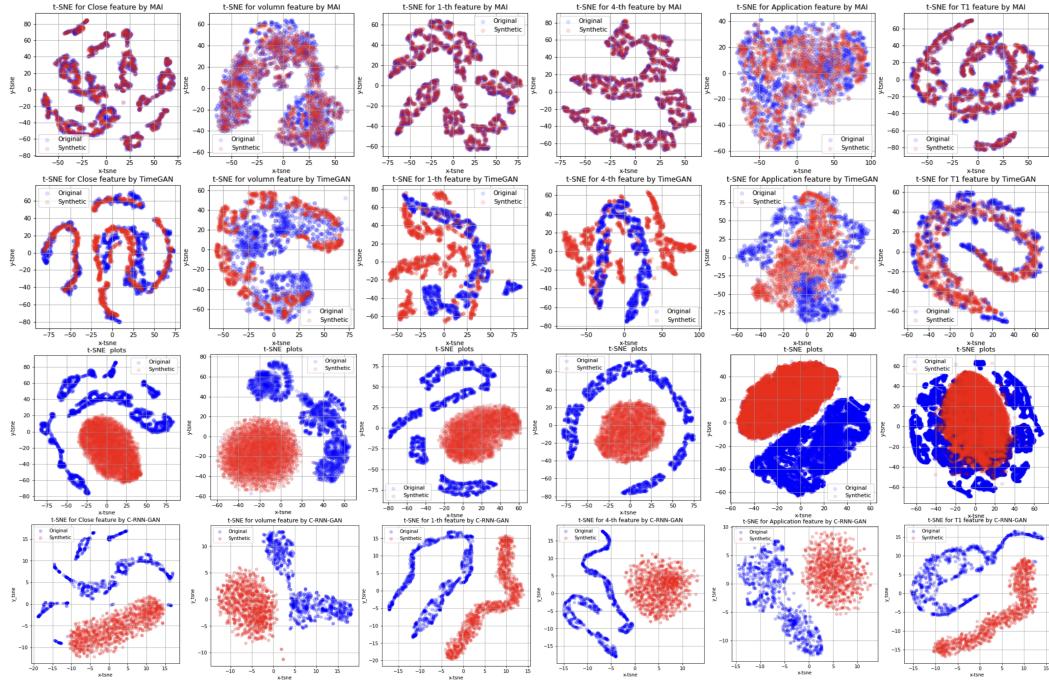
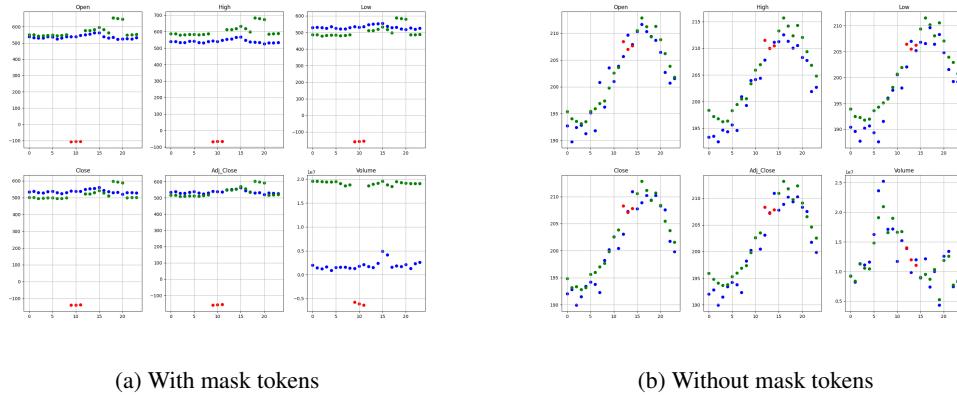


Figure 1: Visualization by t-SNE.

12 We note that RC-GAN and C-RNN-GAN have terrible performance. The manifolds of synthetic data
 13 (red dots) are entirely different from the manifold of original data (blue dots) on row 3 (RCGAN) and
 14 row 4 (C-RNN-GAN). However, synthetic data (red dots in the 1st row) generated by ExtraMAE is
 15 in perfect sync with the original data (blue dots).

16 2 Avoid mask tokens

Figure 2: Illustration for ExtraMAE with mask tokens *v.s* ExtraMAE without mask tokens. We randomly pick up a time series from Stock. Since each time series in Stock has six features, we plot each feature as an independent univariate time series. We first visualize the original time series in blue dots. Then, we mask three consecutive positions in the original time series. ExtraMAE reconstructs the original signal from the partial observation. Then, we plot the reconstructed time series in two colors. For positions visible to ExtraMAE, the reconstructed time series are green dots. For masked positions, we plot the reconstructed time series with red dots. ExtraMAE with mask tokens generates highly discrete reconstruction, but ExtraMAE without mask tokens reconstructs the original signal perfectly.



17 We compare ExtraMAE with mask tokens *v.s* ExtraMAE without mask tokens in Figure 2. Extra-
 18 MAE with mask tokens (figure on the left) generates highly discontinuous reconstruction. Note that
 19 the time series takes a sudden jump in the missing positions (red dots). We guess the reason for such
 20 discontinuity is mask tokens. Therefore, we introduce an extrapolator in our paper to avoid mask
 21 tokens. The ExtraMAE with extrapolator avoids mask tokens (figure on the right) and reconstructs
 22 the original signal perfectly.

23 3 Benchmarks

24 We implements all benchmarks in experiments from their original source codes.

- 25 1. TimeGAN <https://github.com/jsyoon0823/TimeGAN>
- 26 2. RCGAN <https://github.com/ratschlab/RGAN>
- 27 3. C-RNN-GAN <https://github.com/olofmogren/c-rnn-gan>
- 28 4. Mean, Median, Soft, and KNN <https://github.com/iskandr/fancyimpute>
- 29 5. BRITS <https://github.com/caow13/BRITS>

30 4 Consumption

31 Simple algorithms that perform well lie at the core of deep learning. We prune stacked Transformer
 32 blocks which are pretty popular in previous masking-based models. ExtraMAE replaces the heavy
 33 Transformer blocks with lightweight RNNs [5]. The pruning makes ExtraMAE the fastest model for
 34 time series generation.

35 Table 1 lists the number of parameters in each model and the time consumption for 50000 epochs on
 36 Stock. We run all experiments on a GeForce RTX 2080 GPU. Table 1 shows that C-RNN-GAN is
 37 heavy and slow. RC-GAN is the smallest model, but its performance is terrible. TimeGAN achieves
 38 better fidelity and practicality than RC-GAN, but it is much bigger than RC-GAN. TimeGAN also
 39 takes a longer time to train. ExtraMAE is the fastest model with only 2718 trainable parameters.
 40 ExtraMAE only needs 0.24 hours to finish 50000 epochs on Stock. The time ExtraMAE consumed is
 41 4.1%, 9.6%, and 0.3% of the time required by TimeGAN, RC-GAN, and C-RNN-GAN, respectively.

Table 1: Time consumption and the number of parameters on the stock dataset. The unit h in the column Time is *hours*

Model	Params	Time	Pred	Disc
ExtraMAE	2718	0.24h	.037	.071
TimeGAN	48775	5.80h	.053	.256
RCGAN	1218	2.50h	.375	.488
C-RNN-GAN	2523258	72.0h	.086	.500

42 Table 1 also provides the prediction scores (Pred) and discrimination scores (Disc) as references for
 43 the synthetic time series’ practicality and fidelity. Both scores are the lower, the better. We note that
 44 ExtraMAE is consistently and significantly better than other benchmarks in practicality and fidelity.

45 5 TSTR Prediction

46 5.1 TSTR Framework

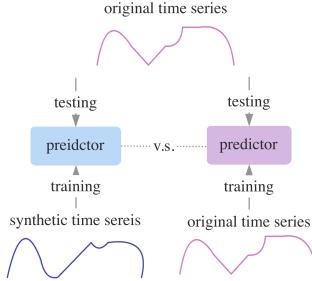


Figure 3: Illustration for TSTR Framework.

47 We consider TSTR on prediction. Usually, we train a predictor (pink on the right) on the original time
 48 series and then test it on the original time series. In TSTR, however, we train the predictor (blue on
 49 the left) on synthetic time series but test it on the original time series. Then, we compare the predictor
 50 trained by synthetic time series (blue on the left) and the predictor trained by the original time series
 51 (pink on the right). The blue predictor is supposed to have a similar performance to the pink one if
 52 the synthetic data is a good substitute for the original data.

53 5.2 Statistics of datasets

54 We provide additional statistics on Stock, Sine, and Energy in Table 2. The Google Stocks dataset is
 55 available online and can be downloaded from LINK. The UCI Appliances Energy Prediction dataset
 56 is also available online and can be downloaded from LINK. Both Stock and Energy are public, and
 57 Sine is simulated. These three datasets are also available in the source code. The datasets contain no
 58 personal information or offensive information.

59 5.3 Hyperparameters

60 We train ExtraMAE 50000 epochs on each dataset since 50000 is the minimal number of epochs that
 61 gives satisfactory results. We also try different batch sizes (i.e., 128, 256, 512, 1024) and found that
 62 the batch size does not significantly influence the model performance. For each dataset, We reserve

Table 2: Statistics of Stock, Sine and Energy. *sequences* is the number of time series in the dataset. We also provide the feature correlations and temporal correlations of the three datasets.

dataset	sequences	features	length	feature corr	temporal variance	temporal corr
Stock	3773	6	24	.8596	.0129	.9902
Sine	10000	5	24	.0117	.3167	.2056
Energy	19711	29	24	.2843	.0444	.8506

63 80% to train ExtraMAE and test it on the left 20%. In Table 3, we list parameters of ExtraMAE
 64 when calculating discrimination scores and prediction scores. Note that L is the length of each time
 65 series, and k is the number of features in each time series. The dataset determines both L and k .

Table 3: Hyperparameters of ExtraMAE on Stock, Sine and Energy

dataset	epochs	batch size	L	k	h	layers	split ratio
Stock	50000	128	24	6	12	3	0.8
Sine	50000	128	24	5	10	3	0.8
Energy	50000	128	24	28	12	3	0.8

66 In ExtraMAE, we implement both encoder and decoder as stacked RNNs followed by a fully
 67 connected layer. The number of layers stacked in the encoder and decoder is the same. In Table 3,
 68 the number of RNN layers stacked is denoted as *layers*. According to Table 3, we implement both
 69 encoder and decoder as 3-layer RNNs followed by a fully connected layer. We find that empirically,
 70 h (the dimension of hidden vector in RNNs) should be twice the number of features k . However, for
 71 time series with many features (e.g., $k = 28$ in Energy), $h = 12$ is enough.

72 6 TSTR classification

73 6.1 Datasets

74 We first describe the three datasets for TSTR classification [3]. Readers may download these three
 75 public datasets from LINK. The datasets contain no personal information or offensive content.

76 **Wafer.** Wafer [6] dataset relates to semiconductor microelectronics fabrication. This dataset consists
 77 of several measurements recorded by sensors in the processing of wafers. The dataset marks time
 78 series data of qualified wafers as normal and labels time-series data produced by defective wafers as
 79 abnormal. The classification problem distinguishes between normal time series and abnormal time
 80 series.

81 **IPD.** IPD [7] stands for Italy Power Demand. This dataset records daily electrical power demand
 82 in Italy for twelve months. The electricity demand patterns in daylight saving time (from April to
 83 September) are quite different from patterns in standard time (from October to March). The classifier
 84 distinguishes time series in daylight saving time from standard time.

85 **Berry.** Berry [8] records sequential spectrographs for fruits. We classify the fruits into two groups,
 86 strawberry and non-strawberry.

87 6.2 Synthetic labels

88 We then discuss how to generate labels for the synthetic time series. In this process, there is no
 89 modification on ExtraMAE, TimeGAN, RC-GAN, and C-RNN-GAN. We take the berry dataset as
 90 an example. Berry has two classes: 0 for non-strawberry and 1 for strawberry. We generate labels in
 91 such ways:

- 92 1. For ExtraMAE: We copy the label of the original time series to the synthetic time series.
 93 For instance, if the label of the original sequence is 1, the label of the synthetic twin will
 94 also be 1. Similarly, we label all synthetic twins of non-strawberry (label 0) with the label 0.

- 95 2. For unsupervised benchmarks (i.e., TimeGAN, RC-GAN, C-RNN-GAN): We train two
 96 independent GANs for sequences of strawberry and sequences belonging to non-strawberry.
 97 We take TimeGAN as an example. First, we collect all original data with label 1 and use this
 98 subset to train a TimeGAN. We call this model TimeGAN_1. Then, we collect all data with
 99 labels 0 and use it to train a new TimeGAN, TimeGAN_0. Finally, we mark synthetic data
 100 produced by TimeGAN_1 with the label 1 and synthetic data generated by TimeGAN_0
 101 with 0.

102 **6.3 hyperparameters**

103 Table 4 provides hyperparameters for TSTR classification. Wafer and Berry contain long time series
 104 of length 152 and 235, but it is enough to use ExtraMAE with hidden vectors of dimension $h = 12$.

Table 4: Hyperparameters of ExtraMAE on Wafer, IPD and Berry

params	epochs	batch size	L	k	h	layers	split ratio
Wafer	50000	128	152	1	12	3	0.8
IPD	50000	128	24	1	12	3	0.8
Berry	50000	128	235	1	12	3	0.8

105 **7 Imputation**

106 We study how well ExtraMAE imputes missing values in the original time series. For each time
 107 series, we randomly mask some patches. We denote the number of mask as n and the mask size as
 108 s. For example, s4n1 indicates 1 mask of length 4. We compute the Mean Squared Errors (MSE)
 109 and Mean Absolute Errors (MAE) in Table 5 6 7 and 8 9 10 respectively. In Table 11 12 13,
 110 we summarize the MSE and MAE for different mask ratios. We repeat one experiment 10 times and
 111 calculate the mean and standard deviations for MSE and MAE.

112 **7.1 MSE of imputation**

Table 5: MSE for Stock

	mean	median	soft	knn	brits	mai
s1n1	0.053212	0.065393	0.000604	0.000559	0.026711	0.000457
s2n1	0.053295	0.065420	0.000718	0.000694	0.024226	0.000695
s3n1	0.052957	0.064955	0.000679	0.000630	0.023400	0.000692
s4n1	0.053323	0.065357	0.000833	0.000804	0.022601	0.000857
s6n1	0.053216	0.065142	0.000875	0.000889	0.021178	0.000865
s8n1	0.053252	0.065368	0.000986	0.001057	0.018279	0.000911
s12n1	0.052941	0.064985	0.145453	0.145453	0.014931	0.001016

Table 6: MSE for Sine

	mean	median	soft	knn	brits	mai
s1n1	0.053627	0.058219	0.000500	7.933572e-07	0.024364	0.001720
s2n1	0.053921	0.058631	0.000652	8.997953e-07	0.021749	0.003865
s3n1	0.054575	0.059366	0.000962	1.100507e-06	0.018255	0.004988
s4n1	0.053867	0.058276	0.001449	1.494250e-06	0.016074	0.001394
s6n1	0.053877	0.058417	0.004048	3.851234e-06	0.013481	0.015626
s8n1	0.053847	0.058221	0.034252	1.924960e-05	0.011687	0.014909
s12n1	0.054224	0.058897	0.346904	3.469041e-01	0.006979	0.005814

Table 7: MSE for Energy

	mean	median	soft	knn	brits	mai
s1n1	0.034796	0.035353	0.007139	0.008325	0.021030	0.009847
s2n1	0.034788	0.035357	0.007325	0.008478	0.019882	0.010233
s3n1	0.034795	0.035358	0.007590	0.008719	0.018985	0.010166
s4n1	0.034730	0.035293	0.007848	0.009051	0.018166	0.010282
s6n1	0.034811	0.035381	0.008588	0.009910	0.016355	0.010543
s8n1	0.034775	0.035345	0.009833	0.010972	0.014538	0.010858
s12n1	0.034744	0.035307	0.250505	0.250505	0.010821	0.010975

113 7.2 MAE of imputation

Table 8: MAE for Stock

	mean	median	soft	knn	brits	mai
s1n1	0.184157	0.168686	0.012415	0.007493	0.097953	0.006553
s2n1	0.184718	0.169127	0.012986	0.008162	0.091715	0.009037
s3n1	0.184153	0.168597	0.013111	0.008488	0.089076	0.010114
s4n1	0.184801	0.169203	0.013939	0.009450	0.084816	0.011661
s6n1	0.184504	0.168790	0.014485	0.011569	0.076509	0.012849
s8n1	0.184479	0.168921	0.015776	0.014755	0.068521	0.014008
s12n1	0.184224	0.168702	0.290512	0.290512	0.052654	0.016312

Table 9: MAE for Sine

	mean	median	soft	knn	brits	mai
s1n1	0.184304	0.179196	0.017113	0.000636	0.108003	0.026997
s2n1	0.184737	0.180204	0.018869	0.000676	0.098634	0.041812
s3n1	0.185969	0.180941	0.021697	0.000738	0.086176	0.049650
s4n1	0.184748	0.179598	0.025256	0.000850	0.077266	0.025592
s6n1	0.184748	0.180200	0.038706	0.001331	0.064739	0.075184
s8n1	0.185373	0.180006	0.130514	0.002675	0.057434	0.080307
s12n1	0.185526	0.180587	0.494050	0.494050	0.037486	0.046360

Table 10: MAE for Energy

	mean	median	soft	knn	brits	mai
s1n1	0.144148	0.141668	0.033155	0.024644	0.094619	0.028022
s2n1	0.144102	0.141647	0.034507	0.025826	0.088664	0.030739
s3n1	0.144077	0.141616	0.036125	0.027580	0.084526	0.033818
s4n1	0.143989	0.141547	0.038019	0.030145	0.080053	0.036042
s6n1	0.144140	0.141694	0.042738	0.036618	0.072286	0.040709
s8n1	0.144097	0.141652	0.049980	0.043422	0.064041	0.045425
s12n1	0.144018	0.141571	0.443258	0.443258	0.047903	0.053531

114 7.3 Summary for imputation

115 In Table 11, 12, 13, we summary the mean and standard error of MSE and MAE on the three datasets:
116 Stock [9], Sine and Energy [10]. We note that ExtraMAE beats all benchmarks with markedly
117 significant improvements. Comparing with the current state-of-the-art BRITS [11], ExtraMAE
118 decrease 97% mean squared error and 86% mean absolute error on the Stock dataset ($97\% \approx$
119 $\frac{|0.0007 - 0.0216|}{0.0216}$, $86\% \approx \frac{|0.0115 - 0.0802|}{0.0802}$). We also present Table 11, 12, 13 in the paper.

Table 11: Averaged MAE and MAE for Stock

method	mse mean	mse std	mae mean	mae std
mean	0.053171	0.000157	0.184434	0.000265
median	0.065231	0.000200	0.168861	0.000232
soft	0.021450	0.054681	0.053318	0.104599
knn	0.021441	0.054685	0.050061	0.106058
brits	0.021618	0.003936	0.080178	0.015581
mai	0.000785	0.000185	0.011505	0.003260

Table 12: Averaged MAE and MAE for Sine

method	mse mean	mse std	mae mean	mae std
mean	0.053991	0.000312	0.185058	0.000579
median	0.058575	0.000428	0.180105	0.000584
soft	0.055538	0.129064	0.106601	0.175513
knn	0.049562	0.131116	0.071565	0.186300
brits	0.016084	0.005977	0.075677	0.024481
mai	0.006902	0.005937	0.049415	0.021431

120 8 Mask ratio

121 8.1 Mask size

122 We first study prediction and discrimination scores under a single mask of different sizes. The length
123 of the original time series is 24. We denote the number of masked positions as the mask size. Table
124 14 and Table 15 shows the results on Stock and Sine respectively. We evaluate the practicality by
125 prediction scores. The lower the prediction score is, the more practical the synthetic data is. Generally,
126 the prediction score grows when the size of the mask increases on Stock and Sine. However, the
127 mask size is not always the smaller; the better. The optimal mask ratio is $50.0\% = \frac{12}{24}$ for Energy,
128 $16.7\% \approx \frac{4}{24}$ for Sine, and $4.2\% \approx \frac{1}{24}$ for Stock. Table 16 shows the results for Energy. Surprisingly,
129 the optimal mask size for Energy is 12 (mask ratio = $\frac{12}{24} = 50.0\%$). For most mask sizes, the
130 prediction scores of ExtraMAE are significantly lower than the scores of the runner-up. We only see
131 exceptions under a few mask sizes (i.e., s12n1 on Stock, s6n1, and s8n1 on Sine).

Table 13: Averaged MAE and MAE for Energy

method	mse mean	mse std	mae mean	mae std
mean	0.034777	0.000030	0.144081	0.000059
median	0.035342	0.000031	0.141628	0.000053
soft	0.042690	0.091643	0.096826	0.152870
knn	0.043709	0.091194	0.090213	0.155820
brits	0.017111	0.003522	0.076013	0.016033
mai	0.010415	0.000400	0.038327	0.008912

132 We evaluate the diversity of synthetic time series by discrimination scores. Low discrimination scores
 133 indicate faithful synthetic data, while high discrimination scores indicate diverse synthetic data. Users
 134 may manage the diversity of synthetic time series generated by ExtraMAE explicitly by changing
 135 the mask ratio. In Table 14 and Table 16, the discrimination scores are consistently lower than the
 136 discrimination score of the runner-up. ExtraMAE achieves great fidelity. In Table 15, ExtraMAE
 137 generates practical (low prediction score) but diverse (high discrimination score) data. By changing
 138 the mask size, we can manage the diversity of synthetic time series directly. Besides, ExtraMAE
 139 succeeds in making a trade-off between diversity and practicality under most mask sizes. Table 14,
 140 Table 15 and Table 16 show results on Stock, Sine, and Energy, respectively. For simplicity, we
 141 denote the mask size as s and the number of masks as n. Therefore, s4n1 stands for an ExtraMAE
 142 trained by under 1 mask of size 4. We repeat each experiment for 10 times for more comprehensive
 143 results and calculate the mean and standard deviation. For example, we train an ExtraMAE on Stock
 144 under 1 mask of size 1 for 10 times and calculate the mean and standard deviation of discrimination
 145 and prediction scores. Finally, we get the mean and standard deviation of ExtraMAE under s1n1 is
 146 0.037010 and 0.110437.

Table 14: Results for different mask sizes on Stock.

method	pred mean	pred std	disc mean	disc std
s1n1	0.037010	0.000131	0.110437	0.116452
s2n1	0.037812	0.000141	0.055457	0.037789
s3n1	0.037264	0.000166	0.114870	0.106279
s4n1	0.040024	0.000248	0.145157	0.078061
s6n1	0.039060	0.000235	0.132196	0.100265
s8n1	0.039412	0.000474	0.102319	0.077281
s12n1	0.060575	0.001133	0.198568	0.019001

Table 15: Results for different mask sizes on Sine.

method	pred mean	pred std	disc mean	disc std
s1n1	0.102111	0.000773	0.273150	0.181229
s2n1	0.103612	0.000542	0.402000	0.163587
s3n1	0.105498	0.001218	0.323100	0.176853
s4n1	0.101471	0.000693	0.113475	0.012641
s6n1	0.136297	0.065392	0.248800	0.084467
s8n1	0.162098	0.001763	0.376700	0.024429
s12n1	0.114557	0.004032	0.370875	0.026211

Table 16: Results for different mask sizes on Energy.

method	pred mean	pred std	disc mean	disc std
s1n1	0.305518	0.010932	0.479292	0.019306
s2n1	0.311318	0.001635	0.440883	0.021442
s3n1	0.296835	0.001658	0.428849	0.014600
s4n1	0.285874	0.002197	0.439120	0.011269
s6n1	0.270136	0.002233	0.402080	0.119129
s8n1	0.278986	0.003768	0.376997	0.142961
s12n1	0.256548	0.000569	0.488372	0.014252

147 8.2 Number of masks

148 Next, we study the prediction and discrimination scores when simultaneously masking several
 149 positions in the time series. The length of the original time series is 24. We try different numbers of
 150 masks of length 1 (The number of masks ranges from 1 to 23). Table 17, Table 18, and Table 19
 151 show the results on Stock, Sine, and Energy respectively.

152 Surprisingly, our ExtraMAE beats the runner up even with an extremely high mask ratio, for example,
 153 $91.7\% \approx \frac{22}{24}$ for Energy, $91.7\% \approx \frac{22}{24}$ for Sine, and $50.0\% = \frac{12}{24}$ for Stock.

Table 17: Different number of masks on Stock

	pred_mean	pred_std	disc_mean	disc_std
s1n1	0.037010	0.000131	0.110437	0.116452
s1n2	0.037449	0.000219	0.128718	0.124518
s1n3	0.037860	0.000249	0.064461	0.026741
s1n4	0.038692	0.000218	0.077285	0.050342
s1n5	0.039692	0.000314	0.163915	0.085648
s1n6	0.041617	0.000472	0.127558	0.053436
s1n7	0.042584	0.000345	0.180491	0.073171
s1n8	0.043500	0.000560	0.135880	0.035202
s1n9	0.045731	0.000491	0.167735	0.019849
s1n10	0.046860	0.001090	0.175784	0.067228
s1n11	0.049972	0.000639	0.170532	0.040340
s1n12	0.051676	0.000901	0.201910	0.016770
s1n13	0.060008	0.000741	0.228922	0.017916
s1n14	0.054770	0.000825	0.189359	0.026818
s1n15	0.060909	0.000959	0.242769	0.048269
s1n16	0.068032	0.000947	0.296794	0.017193
s1n17	0.068534	0.001186	0.268213	0.050064
s1n18	0.074174	0.001060	0.273806	0.022039
s1n19	0.083529	0.001307	0.311664	0.017660
s1n20	0.079422	0.001936	0.276330	0.023975
s1n21	0.083176	0.001841	0.260300	0.033009
s1n22	0.123771	0.002201	0.418349	0.015659
s1n23	0.150588	0.001870	0.366780	0.033819

Table 18: Different number of masks on Sine

	pred_mean	pred_std	disc_mean	disc_std
s1n1	0.102111	0.000773	0.273150	0.181229
s1n2	0.103853	0.001202	0.364625	0.187255
s1n3	0.108489	0.002033	0.402650	0.172492
s1n4	0.125481	0.026787	0.408800	0.160596
s1n5	0.120690	0.011490	0.402550	0.128909
s1n6	0.141998	0.021988	0.452200	0.058036
s1n7	0.143282	0.024395	0.403225	0.097286
s1n8	0.141339	0.025425	0.469525	0.043682
s1n9	0.141593	0.030420	0.461425	0.039935
s1n10	0.128350	0.007239	0.489750	0.015517
s1n11	0.120902	0.017048	0.479800	0.030173
s1n12	0.130095	0.005233	0.485075	0.013444
s1n13	0.129454	0.001641	0.494325	0.006016
s1n14	0.119676	0.001610	0.489625	0.009129
s1n15	0.119031	0.000703	0.464150	0.030883
s1n16	0.122080	0.002530	0.472850	0.035636
s1n17	0.115005	0.000779	0.409375	0.029047
s1n18	0.110826	0.000774	0.446025	0.039793
s1n19	0.107586	0.000390	0.402575	0.073884
s1n20	0.112869	0.013346	0.322475	0.094419
s1n21	0.116344	0.014419	0.369475	0.049660
s1n22	0.106240	0.000992	0.388050	0.088368
s1n23	0.138103	0.043557	0.230275	0.142925

Table 19: Different number of masks on Energy

	pred_mean	pred_std	disc_mean	disc_std
s1n1	0.305518	0.010932	0.479292	0.019306
s1n2	0.276237	0.006469	0.485049	0.010093
s1n3	0.281197	0.004245	0.494674	0.002273
s1n4	0.302781	0.009273	0.495295	0.002193
s1n5	0.316248	0.012123	0.497362	0.002736
s1n6	0.327704	0.013487	0.498085	0.002122
s1n7	0.329568	0.014006	0.499341	0.000467
s1n8	0.347378	0.013903	0.498935	0.001107
s1n9	0.346808	0.013220	0.499696	0.000261
s1n10	0.349471	0.005714	0.499658	0.000227
s1n11	0.349237	0.002552	0.499835	0.000413
s1n12	0.364403	0.002055	0.499391	0.000583
s1n13	0.364816	0.001406	0.495891	0.011536
s1n14	0.362502	0.002799	0.499772	0.000487
s1n15	0.366674	0.001830	0.499366	0.001216
s1n16	0.365429	0.001416	0.499645	0.000408
s1n17	0.367440	0.001106	0.499024	0.002069
s1n18	0.358168	0.002255	0.499822	0.000228
s1n19	0.334868	0.002492	0.499543	0.000834
s1n20	0.322398	0.001711	0.499810	0.000190
s1n21	0.310435	0.001696	0.499848	0.000416
s1n22	0.299452	0.003624	0.499911	0.000081
s1n23	0.307476	0.002883	0.499784	0.000278

154 References

- 155 [1] Laurens Van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of machine*
 156 *learning research*, 9(11), 2008.
- 157 [2] Jinsung Yoon, Daniel Jarrett, and Mihaela Van der Schaar. Time-series generative adversarial
 158 networks. *Advances in Neural Information Processing Systems*, 32, 2019.
- 159 [3] Cristóbal Esteban, Stephanie L Hyland, and Gunnar Rätsch. Real-valued (medical) time series
 160 generation with recurrent conditional gans. *arXiv preprint arXiv:1706.02633*, 2017.
- 161 [4] Olof Mogren. C-rnn-gan: Continuous recurrent neural networks with adversarial training. *arXiv*
 162 *preprint arXiv:1611.09904*, 2016.
- 163 [5] Jeffrey L Elman. Finding structure in time. *Cognitive science*, 14(2):179–211, 1990.
- 164 [6] Wafer dataset. <http://www.timeseriesclassification.com/description.php?Dataset=Wafer>. Donated by: R. Olszewski.
- 165 [7] Ipd dataset. <http://www.timeseriesclassification.com/description.php?Dataset=ItalyPowerDemand>. Donated by: E. Keogh, L.Wi.
- 166 [8] Berry dataset. <http://www.timeseriesclassification.com/description.php?Dataset=Strawberry>. Donated by K. Kemsley and A. Bagnall.
- 167 [9] Stock dataset. https://github.com/jsyoon0823/TimeGAN/blob/master/data/stock_data.csv. Date published: 2021-11-01.
- 168 [10] Energy dataset. <https://archive.ics.uci.edu/ml/datasets/Appiances+energy+prediction>. Date Donated 2017-02-15.
- 169 [11] Wei Cao, Dong Wang, Jian Li, Hao Zhou, Lei Li, and Yitan Li. Brits: Bidirectional recurrent
 170 imputation for time series. *Advances in neural information processing systems*, 31, 2018.